



# Batch/Layer Normalization

2025-02-06

InHyeok Jeong

# Content

- ☐ **Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift**
- ☐ **Layer Normalization**

# Index

- **What is Internal Covariate Shift?**
- **Goal**
- **How to reduce Internal Covariate Shift?**
  - Whiten
  - Batch Normalization
- **Advantages**
- **Accelerating BN Networks**
- **Experiments**
  - MNIST dataset
  - ImageNet classification
  - Ensemble classification
- **Conclusion**


# What is Internal Covariate Shift?

## □ Covariate Shift

- Change in the distributions of layers' inputs
- Require lower learning rates and careful parameter initialization
- Can be applied to whole learning system and its part *e. g.*, sub-network

□  $l = F_2(F_1(u, \Theta_1), \Theta_2), x = F_1(u, \Theta_1) \rightarrow l = F_2(x, \Theta_2)$

□  $\Theta_2 \leftarrow \Theta_2 - \frac{\alpha}{m} \sum_{i=1}^m \frac{\partial F_2(x_i, \Theta_2)}{\partial \Theta_2}$




# What is Internal Covariate Shift?

- Fixed distribution of inputs to a sub-network would have positive consequence for the layer outside the sub-network

- *e.g.*,  $z = g(Wu + b)$ ,  $g(x) = \frac{1}{1+\exp(-x)}$

- As  $|x|$  increases,  $g'(x)$  tends to zero  
→  $x$  moves to the saturated regime




# Goal

- **Reduce Internal Covariate Shift**
  - Make more stable during training
  - Accelerate training

# How to reduce Internal Covariate Shift?

## □ Whiten

- Make inputs have Zero means, unit variances
- Decorrelate



# How to reduce Internal Covariate Shift?

## □ Whiten

### ■ Key consideration

- Gradient descent optimization take into account the normalization
- *e.g.*,  $\hat{x} = x - E[x]$ ,  $x = u + b$ ,  $X = \{x_1, \dots, x_N\}$ ,  $E[x] = \frac{1}{N} \sum_{i=1}^N x_i$
- $b \leftarrow b + \Delta b$ ,  $\Delta b \propto -\frac{\partial l}{\partial \hat{x}}$
- $u + (b + \Delta b) - E[u + (b + \Delta b)] = u + b - E[u + b]$

→ The network *a/ways* produces activations with the desired distribution!



# How to reduce Internal Covariate Shift?

- $x$  : layer input,  $X$  : the set of inputs over the training data set
  - $\hat{x} = \text{Norm}(x, X)$
  - For backpropagation,
    - We need to compute  $\frac{\partial \text{Norm}(x, X)}{\partial x}, \frac{\partial \text{Norm}(x, X)}{\partial X}$
  - Too expensive
    - Require computing  $\text{Cov}[x]$  and derivatives of transforms for backpropagation

# How to reduce Internal Covariate Shift?

## □ Batch Normalization

### ■ Two necessary simplifications

□ Normalize each scalar feature independently

□ For a layer with  $d$ -dimensional input  $x = (x^{(1)} \dots x^{(d)})$

$$\hat{x}^{(k)} = \frac{x^{(k)} - \mathbb{E}[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

□ But simply normalizing each input of layer may change what the layer can represent  
→ The transformation inserted in the network can represent the identity transform

$$y^{(k)} = \gamma^{(k)} \hat{x}^{(k)} + \beta^{(k)}$$

# How to reduce Internal Covariate Shift?

- **Batch Normalization**

- Two necessary simplifications

- *Each mini-batch produces estimates of the mean and variance of each activation*

# How to reduce Internal Covariate Shift?

## □ Batch Normalization

### ■ Accelerate the training of the sub-network

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

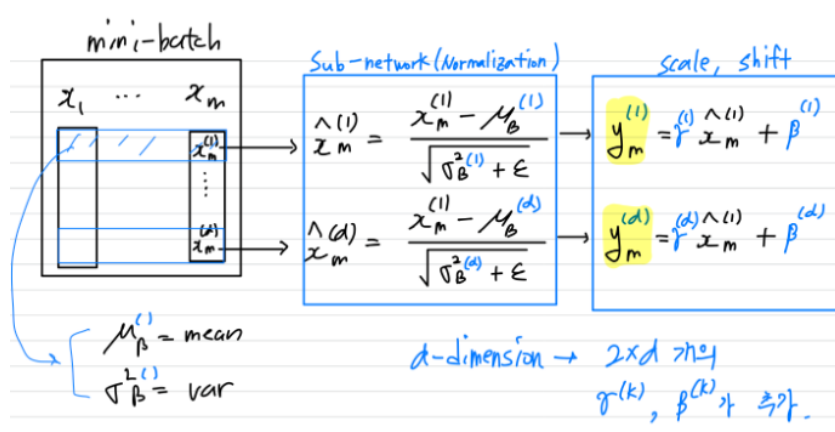
$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.



# How to reduce Internal Covariate Shift?

## □ Batch Normalization

### ■ Can differentiate transformation

$$\frac{\partial \ell}{\partial \hat{x}_i} = \frac{\partial \ell}{\partial y_i} \cdot \gamma$$

$$\frac{\partial \ell}{\partial \sigma_{\mathcal{B}}^2} = \sum_{i=1}^m \frac{\partial \ell}{\partial \hat{x}_i} \cdot (x_i - \mu_{\mathcal{B}}) \cdot \frac{-1}{2} (\sigma_{\mathcal{B}}^2 + \epsilon)^{-3/2}$$

$$\frac{\partial \ell}{\partial \mu_{\mathcal{B}}} = \left( \sum_{i=1}^m \frac{\partial \ell}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \right) + \frac{\partial \ell}{\partial \sigma_{\mathcal{B}}^2} \cdot \frac{\sum_{i=1}^m -2(x_i - \mu_{\mathcal{B}})}{m}$$

$$\frac{\partial \ell}{\partial x_i} = \frac{\partial \ell}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} + \frac{\partial \ell}{\partial \sigma_{\mathcal{B}}^2} \cdot \frac{2(x_i - \mu_{\mathcal{B}})}{m} + \frac{\partial \ell}{\partial \mu_{\mathcal{B}}} \cdot \frac{1}{m}$$

$$\frac{\partial \ell}{\partial \gamma} = \sum_{i=1}^m \frac{\partial \ell}{\partial y_i} \cdot \hat{x}_i$$

$$\frac{\partial \ell}{\partial \beta} = \sum_{i=1}^m \frac{\partial \ell}{\partial y_i}$$

# How to reduce Internal Covariate Shift?

## □ Batch Normalization

- Once the network has been trained, we use the below normalization

$$\hat{x} = \frac{x - \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}}$$

using the population

# How to reduce Internal Covariate Shift?

## □ Batch Normalization

**Input:** Network  $N$  with trainable parameters  $\Theta$ ;  
subset of activations  $\{x^{(k)}\}_{k=1}^K$

**Output:** Batch-normalized network for inference,  $N_{\text{BN}}^{\text{inf}}$

- 1:  $N_{\text{BN}}^{\text{tr}} \leftarrow N$  // Training BN network
- 2: **for**  $k = 1 \dots K$  **do**
- 3:   Add transformation  $y^{(k)} = \text{BN}_{\gamma^{(k)}, \beta^{(k)}}(x^{(k)})$  to  $N_{\text{BN}}^{\text{tr}}$  (Alg. 1)
- 4:   Modify each layer in  $N_{\text{BN}}^{\text{tr}}$  with input  $x^{(k)}$  to take  $y^{(k)}$  instead
- 5: **end for**
- 6: Train  $N_{\text{BN}}^{\text{tr}}$  to optimize the parameters  $\Theta \cup \{\gamma^{(k)}, \beta^{(k)}\}_{k=1}^K$
- 7:  $N_{\text{BN}}^{\text{inf}} \leftarrow N_{\text{BN}}^{\text{tr}}$  // Inference BN network with frozen parameters
- 8: **for**  $k = 1 \dots K$  **do**
- 9:   // For clarity,  $x \equiv x^{(k)}, \gamma \equiv \gamma^{(k)}, \mu_{\mathcal{B}} \equiv \mu_{\mathcal{B}}^{(k)}$ , etc.
- 10:   Process multiple training mini-batches  $\mathcal{B}$ , each of size  $m$ , and average over them:
$$\begin{aligned} \mathbb{E}[x] &\leftarrow \mathbb{E}_{\mathcal{B}}[\mu_{\mathcal{B}}] \\ \text{Var}[x] &\leftarrow \frac{m}{m-1} \mathbb{E}_{\mathcal{B}}[\sigma_{\mathcal{B}}^2] \end{aligned}$$
- 11:   In  $N_{\text{BN}}^{\text{inf}}$ , replace the transform  $y = \text{BN}_{\gamma, \beta}(x)$  with
$$y = \frac{\gamma}{\sqrt{\text{Var}[x] + \epsilon}} \cdot x + \left( \beta - \frac{\gamma \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} \right)$$
- 12: **end for**

**Algorithm 2:** Training a Batch-Normalized Network

# Advantages

- **Make it possible to have a high learning rate**

- It prevents small changes to the parameters from amplifying into larger and suboptimal changes in activations in gradients

- **Make training more resilient to the parameter scale**

- Backpropagation through a layer is unaffected by the scale of its parameters

- $BN(Wu) = BN((aW)u)$

- $\frac{\partial BN((aW)u)}{\partial u} = \frac{\partial BN(Wu)}{\partial u}, \quad \frac{\partial BN((aW)u)}{\partial (aW)} = \frac{1}{a} \cdot \frac{\partial BN(Wu)}{\partial W}$



# Advantages

## □ Regularize the model

- A training example is seen in conjunction with other examples in the mini-batch
- Training network no longer produces deterministic values for a given training example


# Accelerating BN Networks

## □ Increase learning rate


- We can achieve a training speedup from higher learning rates with no ill side effects

## □ Remove or Reduce Dropout

- BN already fulfills some of the same goals as Dropout



(a) Standard Neural Net



(b) After applying dropout.

# Accelerating BN Networks

## □ Reduce the $L_2$ weight regularization

- $L_\lambda(w) = L(w) + \lambda ||w||_2^2$


- When L2 regularization is used with BN, the original regularization effect disappears, leaving only the effect of increasing the learning rate

## □ Shuffle training examples more thoroughly


- Prevent the same examples from always appearing in a mini-batch together

# Experiments


## □ MNIST dataset



(a)



(b) Without BN



(c) With BN

# Experiments

## □ ImageNet classification

■ Apply Batch Normalization to a new variant of the Inception network




Figure 2: Single crop validation accuracy of Inception and its batch-normalized variants, vs. the number of training steps.

Model	Steps to 72.2%	Max accuracy
Inception	$31.0 \cdot 10^6$	72.2%
BN-Baseline	$13.3 \cdot 10^6$	72.7%
BN-x5	$2.1 \cdot 10^6$	73.0%
BN-x30	$2.7 \cdot 10^6$	74.8%
BN-x5-Sigmoid		69.8%

Figure 3: For Inception and the batch-normalized variants, the number of training steps required to reach the maximum accuracy of Inception (72.2%), and the maximum accuracy achieved by the network.

# Experiments

## □ Ensemble classification

Model	Resolution	Crops	Models	Top-1 error	Top-5 error
GoogLeNet ensemble	224	144	7	-	6.67%
Deep Image low-res	256	-	1	-	7.96%
Deep Image high-res	512	-	1	24.88	7.42%
Deep Image ensemble	variable	-	-	-	5.98%
BN-Inception single crop	224	1	1	25.2%	7.82%
BN-Inception multicrop	224	144	1	21.99%	5.82%
BN-Inception ensemble	224	144	6	20.1%	<b>4.9%*</b>

Figure 4: *Batch-Normalized Inception comparison with previous state of the art on the provided validation set comprising 50000 images. \*BN-Inception ensemble has reached 4.82% top-5 error on the 100000 images of the test set of the ImageNet as reported by the test server.*

# Conclusion

- Batch Normalization reduces Internal Covariate Shift, accelerating training and improving stability
- It enhances generalization, often eliminating the need for Dropout
- By further increasing the learning rates, removing Dropout, and applying other modifications afforded by Batch Normalization

# Index

## ☐ Disadvantages of Batch Normalization

## ☐ Layer Normalization

- Vs. Batch Normalization
- Similarity
- Difference

## ☐ Experiments

- Order embeddings of images and language
- Handwriting sequence generation
- Permutation invariant MNIST
- Convolutional Networks

## ☐ Conclusion




# Disadvantages of Batch Normalization

- ❑ Dependent on the mini-batch size
- ❑ Not obvious how to apply it to recurrent neural networks

# Layer Normalization

## □ Vs. Batch Normalization

Batch Normalization




$$\mu_i^l = \mathbb{E}_{\mathbf{x} \sim P(\mathbf{x})} [a_i^l] \quad \sigma_i^l = \sqrt{\mathbb{E}_{\mathbf{x} \sim P(\mathbf{x})} [(a_i^l - \mu_i^l)^2]}$$

# Layer Normalization

## □ Vs. Batch Normalization

Layer Normalization



$$\mu^l = \frac{1}{H} \sum_{i=1}^H a_i^l \quad \sigma^l = \sqrt{\frac{1}{H} \sum_{i=1}^H (a_i^l - \mu^l)^2}$$

# Layer Normalization

## ☐ Similarity

- Give each neuron its own adaptive bias and gain

## ☐ Difference

- Perform exactly the same computation at training and test times

# Experiments

## □ Order embeddings of images and language




Figure 1: Recall@K curves using order-embeddings with and without layer normalization.

MSCOCO								
Model	Caption Retrieval				Image Retrieval			
	R@1	R@5	R@10	Mean $r$	R@1	R@5	R@10	Mean $r$
Sym [Vendrov et al., 2016]	45.4		88.7	5.8	36.3		85.8	9.0
OE [Vendrov et al., 2016]	46.7		88.9	5.7	37.9		85.9	8.1
OE (ours)	46.6	79.3	89.1	5.2	37.8	73.6	85.7	7.9
OE + LN	<b>48.5</b>	<b>80.6</b>	<b>89.8</b>	<b>5.1</b>	<b>38.9</b>	<b>74.3</b>	<b>86.3</b>	<b>7.6</b>

Table 2: Average results across 5 test splits for caption and image retrieval. **R@K** is Recall@K (high is good). **Mean  $r$**  is the mean rank (low is good). Sym corresponds to the symmetric baseline while OE indicates order-embeddings.

# Experiments

## □ Handwriting sequence generation




Figure 5: Handwriting sequence generation model negative log likelihood with and without layer normalization. The models are trained with mini-batch size of 8 and sequence length of 500,

# Experiments

## □ Permutation invariant MNIST

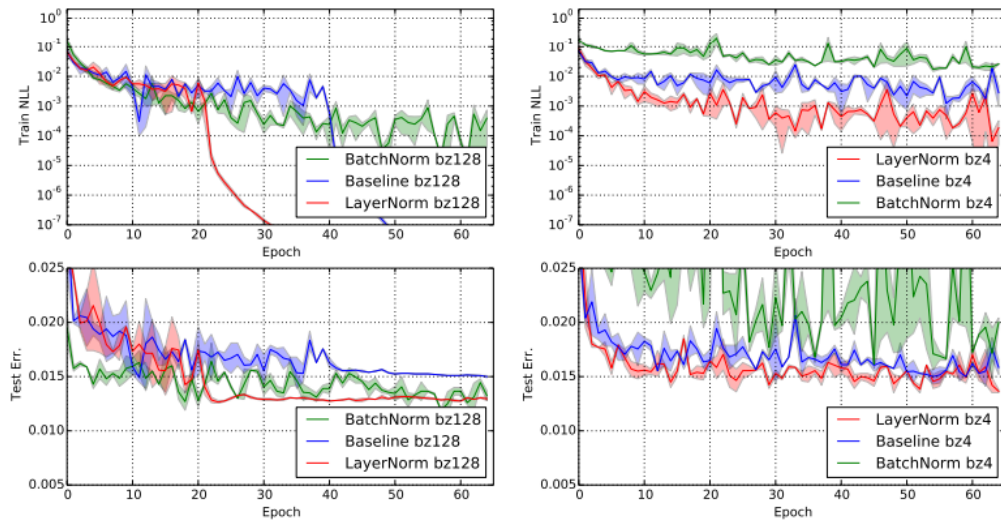


Figure 6: Permutation invariant MNIST 784-1000-1000-10 model negative log likelihood and test error with layer normalization and batch normalization. (Left) The models are trained with batch-size of 128. (Right) The models are trained with batch-size of 4.

# Experiments

## □ Convolutional Networks

- The large number of the hidden units whose receptive fields lie near the boundary of the image are rarely activated
- Have very different statistics from the rest of the hidden units within the same layer
- BN outperforms LN



# Conclusion

- ☐ **Not dependent on mini-batch size compared to Batch Normalization**
- ☐ **Applicable to RNN**
- ☐ **Further research is needed to make layer normalization work well in CNN**