



Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling

2025-07-08

ICLR 2014

Junyoung Chung, Caglar Gulcehre, KyungHyun Cho, Yoshua Bengio

☐ Introduction

- What is a Recurrent Neural Network?
- Types of a Recurrent Neural Network
- Long-Term Dependency Problem

☐ Gated Recurrent Neural Networks


- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

☐ Experiments

☐ Conclusion

What is a Recurrent Neural Network?

- An extension of a conventional feedforward neural network, which is able to handle a variable-length sequence input
 - Text, Music, Stock price, ...



automated data mining survey
responses com ter transcripts
qualatativ root cause
classification insights
ad-hoc analysis product
reviews ser it vol of the
customer dashboards consumer
trends ad-hoc analysis early warning

What is a Recurrent Neural Network?

□ $x = (x_1, x_2, \dots, x_T)$

■ Input sequence


□ $h_t = g(Ux_t + Wh_{t-1})$

■ Recurrent hidden state (state vector)

■ Contain information about the history of all the past elements of the sequence

□ $y = (y_1, y_2, \dots, y_T)$



■ Output (optional)



Types of Recurrent Neural Networks

□ Many to One


■ *e. g.*, Emotional analysis



Types of Recurrent Neural Networks

□ Many to Many

■ *e. g.*, Named Entity Recognition



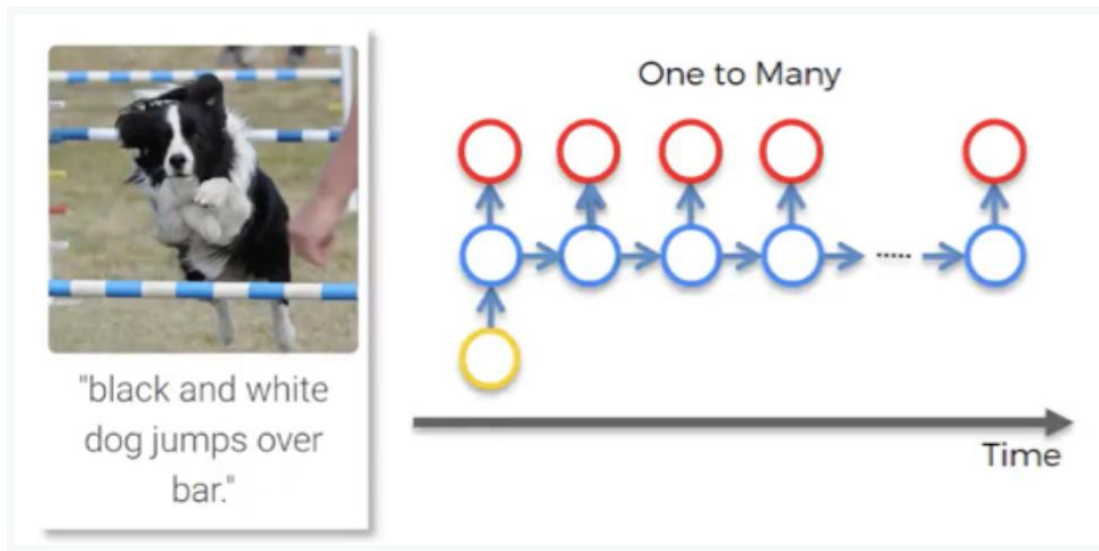
Person p Loc l Org o Event e Date d Other z

Barack Hussein Obama II * (born August 4, 1961 *) is an American * attorney and politician who served as the 44th President of the United States * from January 20, 2009 *, to January 20, 2017 *. A member of the Democratic Party *, he was the first African American * to serve as president. He was previously a United States Senator * from Illinois * and a member of the Illinois State Senate *.

Types of Recurrent Neural Networks

□ One to Many

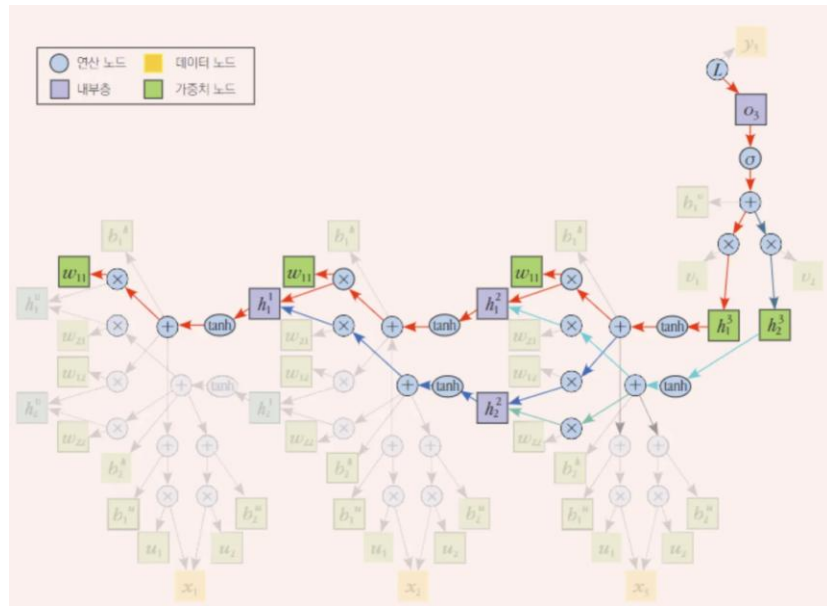
■ *e. g.*, Image Captioning



Long-Term Dependency Problem

□ Difficult to train RNNs to capture Long-Term Dependencies

- The gradients tend to either vanish or explode




Long-Term Dependency Problem

□ Solution

■ Devise a better learning algorithm

- e. g., Clip gradients above a given threshold




Long-Term Dependency Problem


□ Solution

■ Use gate units

- Perform well in tasks that require capturing long-term dependencies



(a) Long Short-Term Memory



(b) Gated Recurrent Unit

Long Short-Term Memory (LSTM)

□ Memory cell

- Explicitly store information up to this time

- Updated by

 - 1) Partially forgetting the existing memory

 - 2) Adding a new memory content


- $c_t^j = f_t^j c_{t-1}^j + i_t^j \tilde{c}_t^j$

- $\tilde{c}_t^j = \tanh(W_c x_t + U_c h_{t-1})^j$

□ Output

- The activation of the LSTM unit

- $h_t^j = o_t^j \tanh(c_t^j)$



Long Short-Term Memory (LSTM)

□ Forget gate

- Modulate **the extent to which the existing memory is forgotten**

- $f_t^j = \sigma(W_f x_t + U_f h_{t-1} + V_f c_{t-1})^j$

□ Input gate

- Modulate **the degree to which the new memory content is added to the memory cell**

- $i_t^j = \sigma(W_i x_t + U_i h_{t-1} + V_i c_{t-1})^j$

□ Output gate

- Modulate **the amount of memory content exposure**

- $o_t^j = \sigma(W_o x_t + U_o h_{t-1} + V_o c_{t-1})^j$

Gated Recurrent Unit (GRU)

□ Have no separate memory cells


□ Linear interpolation

■ The previous activation h_{t-1}^j

■ The candidate activation \tilde{h}_t^j

■ $h_t^j = (1 - z_t^j)h_{t-1}^j + z_t^j\tilde{h}_t^j$

■ $\tilde{h}_t^j = \tanh(Wx_t + U(r_t \odot h_{t-1}))^j$



Gated Recurrent Unit (GRU)

□ Update gate

- Decide **how much the unit updates its activation, or content**

- $z_t^j = \sigma(W_z x_t + U_z h_{t-1})^j$

□ Reset gate

- Determine **the percentage of oblivion to past information** based on

- 1) Information of the current timestamp
- 2) Past information

- $r_t^j = \sigma(W_r x_t + U_r h_{t-1})^j$

☐ Similarities

- Keep the existing content and add the new content on top of it

☐ Differences

- Control of the amount of information to the next timestamp
- Oblivion and renewal of information

☐ Advantages

- Easy for each unit to remember the existence of a specific feature in the input stream
- Reduce the difficulty due to vanishing gradients via the addition

☐ Tasks

■ Sequence modeling

- ☐ Aim at learning a probability distribution over sequences

→ *Maximize the log-likelihood of a model!*

- ☐ $\max_{\theta} \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \log p(x_t^n | x_1^n, \dots, x_{t-1}^n; \theta)$

☐ Datasets

■ Music Datasets

- ☐ Polyphonic music modeling
- ☐ Nottingham, JSB Chorales, MuseData, Piano-midi

■ Ubisoft Datasets

- ☐ Speech signal modeling
- ☐ Ubisoft A, Ubisoft B

□ Models

- Make models have approximately the same number of parameters

Unit	# of Units	# of Parameters
Polyphonic music modeling		
LSTM	36	$\approx 19.8 \times 10^3$
GRU	46	$\approx 20.2 \times 10^3$
tanh	100	$\approx 20.1 \times 10^3$
Speech signal modeling		
LSTM	195	$\approx 169.1 \times 10^3$
GRU	227	$\approx 168.9 \times 10^3$
tanh	400	$\approx 168.4 \times 10^3$

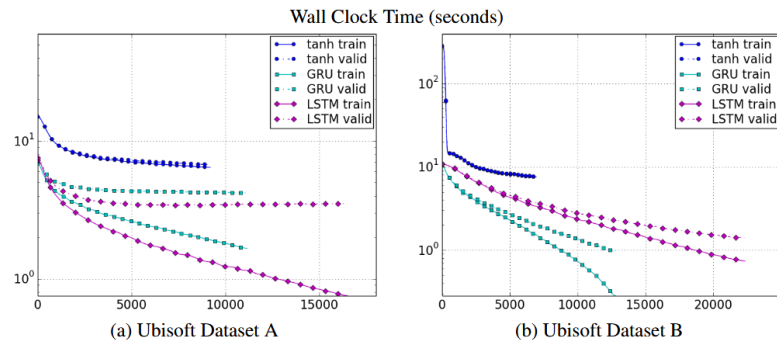

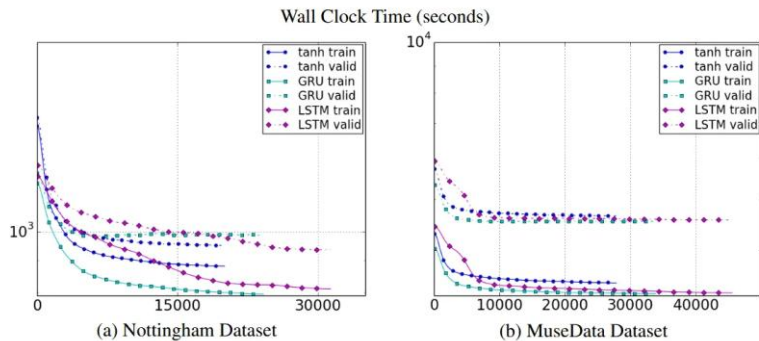
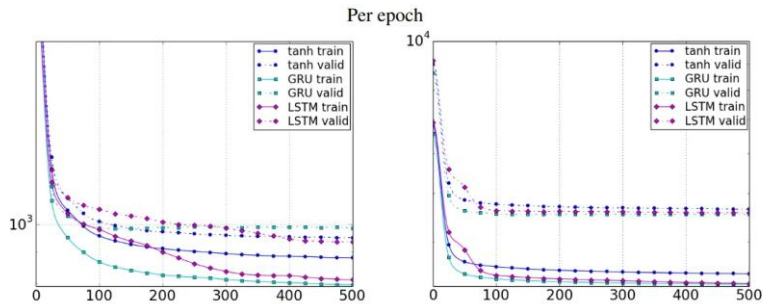
Table 1: The sizes of the models tested in the experiments.

□ The average negative log-probabilities of the training and test sets

■ GRU-RNN outperformed all the others on all the datasets except for the Nottingham

			tanh	GRU	LSTM
Music Datasets	Nottingham	train	3.22	2.79	3.08
		test	3.13	3.23	3.20
	JSB Chorales	train	8.82	6.94	8.15
		test	9.10	8.54	8.67
	MuseData	train	5.64	5.06	5.18
		test	6.23	5.99	6.23
	Piano-midi	train	5.64	4.93	6.49
		test	9.03	8.82	9.03
Ubisoft Datasets	Ubisoft dataset A	train	6.29	2.31	1.44
		test	6.44	3.59	2.70
	Ubisoft dataset B	train	7.61	0.38	0.80
		test	7.62	0.88	1.26

□ Learning curves of the best validation runs



☐ What is a Recurrent Neural Network?

- Address sequence data
- Many to One, Many to Many, One to Many
- Long-Term Dependency Problem

☐ Gated Recurrent Neural Network

- Use several gates to catch long term dependencies

☐ Experiments

- Gated Neural Networks outperform traditional RNN